

k8slens

# The Most Common Kubernetes Errors and How to Fix Them



Flavius Dinu

Follow

11 min read · Sep 1, 2025

67



The infographic features a dark background with a subtle geometric pattern of orange and blue lines. In the center, there's a white 'LENS' logo with a camera shutter icon. To the left, the text 'Fixing the most common K8s errors' is displayed in large white font, followed by 'With Lens K8s IDE' in blue font. To the right, seven blue rounded rectangles list common Kubernetes errors: 'CrashloopBackoff', 'ImagePullBackoff', 'CreateContainerError', 'OOMKilled', 'NodeNotReady', 'PodInitializingError', and 'PodEviction'. Each error name is preceded by a small white icon.

Managing Kubernetes (K8s) at scale can easily transform into a cumbersome process. There are many issues that can arise, which can be hard to solve, especially if you don't know where to look, or you haven't faced them before.

Even if you've seen these errors in the past, solving them can have different solutions, depending on their source: pod scheduling, container runtime issues, misconfigured manifests, network policies, and others.

In this article, we will take a look at what are the most common K8s errors, and how easy it is to fix them with Lens, and Lens Prism.

## What are the most common Kubernetes Errors?

Let's explore an overview of what are the most common errors, and their most common solutions. You will see that there are many different solutions you will need to try for these errors, thus making the debugging process hard.

There is however, a formula that can be always used for debugging K8s issues when they are related to pods. You should always try to:

- Use `kubectl describe` to understand what is happening with your pod
- Check the pod logs
- Verify the pod manifest

TL;DR? Check out the video instead:

## 8 Kubernetes Errors That Will Destroy Your Deployments (And How To Fix ...)



### 1. CrashLoopBackoff

**CrashLoopBackoff** means that a pod is repeatedly crashing during startup and it is being restarted by the `kubelet`. There are many reasons for why this can happen, and some of the most common ones include:

- Misconfigured command or entrypoint
- Resource limits for CPU or Memory are too low
- Pods may be trying to use non-existent resources
- Missing environment variables or config files
- Application error causing the container to exit
- Containers might fail on missing file permissions
- Incompatible base images
- Database readiness is not ensured before the application starts

This is how a pod with this status would look like in the CLI:

```
kubectl get pods NAME READY STATUS RESTARTS AGE
instant-crash-pod 0/1 CrashLoopBackOff 4 (65s ago) 2m44s
```

There are many other reasons for why **CrashLoopBackoff** may occur, so in order to fix it, you will need to investigate a couple of things.

The first step you will need to do is to run `kubectl describe` on the failed pod to inspect the events that have happened during the creation of the pod. Depending if the pod is deployed as part of a deployment/statefulset/replicaset/daemonset/standalone you will need to run the appropriate `kubectl describe` command.

Example output for a `describe` command for a pod that is in the **CrashLoopBackoff** state:

Events:

Type	Reason	Age	From	Message
Normal	Scheduled	3m43s	default-scheduler	Successfully assigned default/instant-crash-pod to lens-kind-control-plane
Normal	Pulled	3m43s	kubelet	Successfully pulled image "busybox" in 935ms (935ms including waiting).
Normal	Pulled	3m41s	kubelet	Successfully pulled image "busybox" in 1.067s (1.067s including waiting).
Normal	Pulled	3m24s	kubelet	Successfully pulled image "busybox" in 965ms (965ms including waiting).
Normal	Pulled	2m59s	kubelet	Successfully pulled image "busybox" in 1.069s (1.069s including waiting).

Normal Pulled 2m5s kubelet Successfully pulled image "busybox" in 1.171s (1.171s including waiting).

Normal Pulling 37s (x6 over 3m44s) kubelet Pulling image "busybox"

Normal Created 35s (x6 over 3m43s) kubelet Created container: crash

Normal Started 35s (x6 over 3m43s) kubelet Started container crash

Normal Pulled 35s kubelet Successfully pulled image "busybox" in 1.129s (1.129s including waiting).

Warning BackOff 11s (x18 over 3m40s) kubelet Back-off restarting failed container crash in pod instant-crash-pod\_default

After you get a sense of what is happening using `kubectl describe`, it is now time to check the pod logs and try to understand why the pod is failing. If your pod has dependencies, make sure you use `initContainers` to delay the creation of the main one, until dependencies are up. At the same time, if you see errors related to liveness and readiness probes, modify their timing or temporarily disable them to verify if they're the issue.

Your pod might be also entering the `CrashLoopBackoff` state if there are issues with memory limits, so ensure that your limits are big enough for the pod to be created.

## 2. ErrImagePull/ImagePullBackOff

`ErrImagePull` is one of the most common Kubernetes errors you will encounter, and this is happening when your cluster cannot pull a container image required by your pods.

This error can happen for several reasons such as:

- Using an incorrect image name or tag

- Private registry access denied
- Image doesn't exist anymore
- Rate limits
- Pod doesn't have permission to access the image

Kubernetes will initially show the **ErrImagePull**, but after retrying multiple times, it will end up with the **ImagePullBackOff** error. After each image pull retry, the delay between scheduling a new download attempt increases.

To fix this issue, you should first check your pod specification and see if it provides the correct image/registry. You can use `docker pull image_name`, to verify if the image can be pulled.

In case you are using private registries, make sure you create the necessary pull secret for that particular registry, and ensure this is attached to your pod in the `imagePullSecrets`.

```
imagePullSecrets: - name: registry_secret_name
```

### 3. CreateContainerError / CreateContainerConfigError

Both **CreateContainerError** and **CreateContainerConfigError**, are common error messages you might encounter when you are deploying resources inside your Kubernetes clusters.

These errors happen before a container starts, but they have different root causes.

**CreateContainerError** typically occurs when:

- There is a bad command or entrypoint, or they are missing altogether
- There is naming conflict
- The container tries to run an executable file without execute permissions
- A file or binary expected in the image doesn't exist
- There are inaccessible storage volumes

**CreateContainerConfigError** typically occurs when:

- There is a missing ConfigMap or a Secret
- You reference a key that doesn't exist in a Secret or ConfigMap

When facing the **CreateContainerError**, you should check the container's events by using the command, understand what the *Failed* event is, and troubleshoot the issue.

On the other hand, when you are facing the **CreateContainerConfigError**, after following the same steps, you will usually see that your pod couldn't find a particular ConfigMap, Secret, or a key in the Configmap or Secret. As soon as you identify the issue, you should modify the Pod manifest and reapply it.

## 4. OOMKilled

The **OOMKilled** error happens when a container is killed because it was using more memory than it was allowed.

Here are some common causes for why this happens:

- There are no memory limits set, or they are too low
- There is an in-app memory spike or leak
- JVM app over-allocates heap
- Large buffers or in-memory caches

After using `kubectl describe` to view all the relevant events from your pod, you should always use the `kubectl top pod` command to understand the pod's memory usage. In case there are no spikes or memory leaks, it is safe to assume that your app consumes more memory than it has allocated, so you need to increase the limits:

```
memory: "128Mi"    memory: "256Mi"
```

You should also ensure that your node has enough resources to handle your workloads, and in some cases, you can also refactor your application to reduce memory consumption.

## 5. NodeNotReady

Nodes can be in different states such as Ready (when it is healthy and capable of accepting new pods), Not Ready (when it has issues, and it is not functioning properly), and Unknown (usually this happens when there are communication issues with the node).

The **NodeNotReady** error typically means that your node cannot run new pods.

Some of the common causes for why this happens include:

- Resource pressure such as low disk space or memory exhaustion
- Issues with the kubelet
- Issues with kube-proxy
- Network problems

To understand what is happening you should check the node status and conditions:

```
kubectl describe node <node-name>
```

Another thing you should do is connect to the node, and see the `kubelet` logs, as this will give you more insights:

```
journalctl -u kubelet
```

In addition to this, you should try to verify there are no network problems, and that kube-proxy is running properly.

## 6. Unschedulable Pods: FailedScheduling

This error means that a pod is stuck in a *Pending* state, because Kubernetes cannot assign it to any node.

Here are some of the most common causes for why this happens:

- Insufficient cluster resources
- Node taints without matching tolerations
- Affinity or node selector rules are overly restrictive

- Nodes are cordoned or drained

To troubleshoot you should first use the *describe* command on the Pod to see if there are any scheduling errors (look for FailedScheduling) events.

Some other actions you can do include:

- Using *describe* commands on the node to check resource usage
- Review taints and tolerations
- Adjust affinity rules
- Uncordon your node

## 7. Readiness/Liveness Probe Failures

Configuring probes to test your application is a best practice when it comes to using Kubernetes. The liveness and readiness probes are very important, but when you set them up, you will most likely start facing a new set of errors:

- Your applications might not be ready during the probe windows
- You have configured a wrong path or port

To fix errors related to these probes, you should verify your manifests and ensure your applications have enough time to be ready before the probe windows, and also double check your paths and your ports.

You can also port-forward your applications and test them manually to ensure everything is working properly.

## 8. PodInitializing Error

If your init container fails or crashes, your main container will not start and the pod will be stuck in the **PodInitialization** state.

The init container can fail for many reasons, from the list described above, so to ensure a smooth debugging process, you should first try to identify what is the case in which your init container is, and follow the debugging process accordingly.


Your troubleshooting checklist should like this:

- Describe the pod and review the init container events
- Try to view the init logs if that's possible
- Understand what the error is, and start implementing one of the fixes

## How can Lens The Kubernetes IDE help you solve these issues?

Lens K8s IDE equips you with everything you need to connect, observe, and manage your Kubernetes clusters. It is a developer product meant to complement `kubectl` and observability platforms.

In Lens K8s IDE, you can see at a glance what are the pods that have issues by checking for the yellow triangle in the *Pods view*.




The screenshot shows the Lens K8s IDE interface with the 'Pods' view open. The title bar says 'Pods - kind-kind'. The search bar contains 'instant-crash-pod'. Below the search bar, there are dropdown menus for 'Namespace' (set to 'default') and 'Container' (set to 'Search Pods...'). A download icon and a count of '6 items' are also present. The main table has columns: Name, Namespace, Containers, CPU, Memory, Restarts, Controlled By, Node, QoS, and Age. There are three rows of data:

Name	Namespace	Containers	CPU	Memory	Restarts	Controlled By	Node	QoS	Age
instant-crash-pod	default	default	N/A	N/A	33	lens-kind	BestEffort	147m	
jenkins-1754476299-fc8c98954-pl2gk	default	default	N/A	N/A	0	ReplicaSet	lens-kind	Burstable	27h
nginx-deployment-55547d4d99-hr6pd	default	default	N/A	N/A	0	ReplicaSet	lens-kind	Burstable	30h

Open in app ↗



understand what is the issue with your pod:



Summary	Count	Age
(combined from similar events): Successfully pulled image "busybox" in 984ms (28 events)	28	6m3s
Back-off restarting failed container crash in pod instant-crash-pod_default(292k events)	808	175m
Pulling image "busybox"	40	175m

Lens K8s IDE also gives you an easy way to connect to the pod shell or attach to the pod, and you can also get the logs without needing to type any `kubectl` commands.

```
Displaying logs from Namespace: default for Pod: nginx-deployment-586bfb775f-szqs8. Logs from 07/08/2025, 17:53:48 EEST

/docker-entrypoint.sh: /docker-entrypoint.d/ is not empty, will attempt to perform configuration
/docker-entrypoint.sh: Looking for shell scripts in /docker-entrypoint.d/
/docker-entrypoint.sh: Launching /docker-entrypoint.d/10-listen-on-ipv6-by-default.sh
10-listen-on-ipv6-by-default.sh: info: Getting the checksum of /etc/nginx/conf.d/default.conf
10-listen-on-ipv6-by-default.sh: info: Enabled listen on IPv6 in /etc/nginx/conf.d/default.conf
/docker-entrypoint.sh: Launching /docker-entrypoint.d/20-envsubst-on-templates.sh
/docker-entrypoint.sh: Launching /docker-entrypoint.d/30-tune-worker-processes.sh
/docker-entrypoint.sh: Configuration complete; ready for start up
2025/08/07 14:53:48 [emerg] 1#1: "server" directive is not allowed here in /etc/nginx/nginx.conf:1
nginx: [emerg] "server" directive is not allowed here in /etc/nginx/nginx.conf:1
```

For node related issues, you can easily see all your nodes, and details about their health, resource allocations, how many pods they are hosting and more.

Capacity					
CPU	Memory	Ephemeral Storage	Hugepages-1Gi	Hugepages-2Mi	Pods
5	7.7GiB	58.4GiB	0	0	110
Allocatable					
CPU	Memory	Ephemeral Storage	Hugepages-1Gi	Hugepages-2Mi	Pods
5	7.7GiB	58.4GiB	0	0	110

If you want to see what other things you can do with Lens K8s IDE, check out [this article](#).

Even if this speeds up your debugging process a lot when compared to plain kubectl, we wanted to take it up a notch, and that's why we've introduced Lens Prism.

## What is Lens Prism and how can it help with the most common Kubernetes Errors?

Lens Prism is an AI-Powered Kubernetes Copilot accessible directly from Lens K8s IDE. It connects to OpenAI compatible LLM and it can help you understand what is happening inside of your Kubernetes clusters, making the debugging process easier.

For this example I will use an OpenAI API key, and leverage GPT 4.1 to get solutions for the issues inside my cluster.

Let's see it in action:

<input type="checkbox"/>	Name	<input type="checkbox"/>	Namespace	Containers
<input type="checkbox"/>	hello-world	<span style="color: yellow;">!</span>	default	<span style="background-color: orange;">■</span>
<input type="checkbox"/>	instant-crash-pod	<span style="color: yellow;">!</span>	default	<span style="background-color: orange;">■</span>
<input type="checkbox"/>	jenkins-1754476299-fc8c98954-pl2gk		default	<span style="background-color: green;">■</span>
<input type="checkbox"/>	nginx-deployment-55547d4d99-kxnks		default	<span style="background-color: green;">■</span>
<input type="checkbox"/>	nginx-deployment-586bfb775f-n27lf	<span style="color: yellow;">!</span>	default	<span style="background-color: orange;">■</span>
<input type="checkbox"/>	nginx-deployment-586bfb775f-n28z8	<span style="color: yellow;">!</span>	default	<span style="background-color: orange;">■</span>
<input type="checkbox"/>	nginx-deployment-586bfb775f-szqs8	<span style="color: yellow;">!</span>	default	<span style="background-color: orange;">■</span>

In the above example, you can see I have many pods that are in an unhealthy state. I will now go into my Lens Prism chat, and ask what is happening with my pods:

What is happening with my pods?

Several pods in your cluster are experiencing issues:

- Multiple pods have extremely high restart counts (e.g., instant-crash-pod: 533, nginx-deployment-586bfb775f-n27lf: 541, nginx-deployment-586bfb775f-szqs8: 541, nginx-deployment-586bfb775f-n28z8: 328). This indicates repeated crashes or failures.
- hello-world pod is stuck in Pending state and has never started.
- Some pods are Running but not Ready, which means their containers are not passing readiness checks.

These are clear signs of instability or misconfiguration. Would you like a deeper analysis of a specific pod, or should I investigate the most critical issues cluster-wide (e.g., high-restart pods and Pending pods)? Please specify your focus.

For a full pod list and details, verify in Lens.

We can easily see a list of all the pods that have issues in the chat, but now lets dig even deeper into their root causes and how to solve them:

Can you investigate all pods that are not running, show me the root causes, and provide fixes?

Here are the root causes and fixes for all pods that are not running:

1. hello-world (Pending, ImagePullBackOff)
  - Root Cause: The image name busybozx is invalid (should be busybox).
  - Fix: Update the pod spec to use the correct image name busybox.
2. instant-crash-pod (CrashLoopBackOff)
  - Root Cause: The container command is false, which always exits with code 1.
  - Fix: Change the command to a valid one (e.g., sleep 3600 or sh -c "echo Hello World").
3. nginx-deployment-586fb775f-n27lf, nginx-deployment-586fb775f-szqs8, nginx-deployment-586fb775f-n28z8 (CrashLoopBackOff)
  - Root Cause: The nginx container is crashing with exit code 1. Likely due to a misconfigured nginx.conf from the nginx-config ConfigMap.
  - Fix: Check the nginx-config ConfigMap for syntax errors or invalid configuration.

Suggested kubectl patch commands (run manually):

- For hello-world:

Copy

```
kubectl patch pod hello-world -n default --type='json' -p='[{"op":"replace","path":"/spec/containers/0/image","value":"busybox"}]'
```

- For instant-crash-pod:

Copy

```
kubectl patch pod instant-crash-pod -n default --type='json' -p='[{"op":"replace","path":"/spec/containers/0/command","value":["sleep","3600"]}]'
```

- For nginx pods: Review and fix the nginx-config ConfigMap, then restart the pods.

Please verify these findings and patches in Lens before applying. Let me know if you want help inspecting the nginx-config ConfigMap or need further details.

We get a list of potential fixes for all of our containers. Let's apply the patches mentioned above and see if this solves are issues for the first two pods:

```
kubectl patch pod hello-world -n default --type='json' -
p='[{"op":"replace","path":"/spec/containers/0/image", "value":"busybox"}]'

pod/hello-world patched
```

The *hello-world* pod was patched successfully and the issue has disappeared.

For the second case related to the *instant-crash-pod*, we are trying to patch a command, and this is forbidden in Kubernetes:

```
kubectl patch pod instant-crash-pod -n default --type='json'...
```

The Pod "instant-crash-pod" is invalid: spec: Forbidden: pod updates may not change fields other than ...

But if we go back to what Lens Prism told us initially, we understand that the command we are using is invalid one, so what we can do is modify the manifest, delete the pod and recreate it. As soon as we do that we can see that it is working as expected.

Now let's deep dive into what is happening with nginx. For that I will ask a follow up question:

The screenshot shows a dark-themed interface with a message at the top right: "Can you tell me what is happening exactly with the nginx-config ConfigMap?". Below it, a message states: "The nginx-config ConfigMap contains an invalid nginx.conf:". A bulleted list follows: • The server block is missing a closing }. • The location block is missing a closing }. • The return statement's string is split across lines, which is not valid in nginx syntax. Another message below says: "This malformed configuration causes nginx to fail on startup, resulting in CrashLoopBackOff for all nginx pods using this ConfigMap.". A "Fix" section provides an example of a valid configuration:

```
nginx
server {
    listen 80;
    location / {
        return 200 'Hello from the Lens tutorial!\n';
        add_header Content-Type text/plain;
    }
}
```

At the bottom, a note says: "Update the ConfigMap and restart the pods. Please verify and apply the fix in Lens. Let me know if you need the exact kubectl command to patch the ConfigMap."

So it's clear that our ConfigMap has some issues, let's fix them as Lens Prism told us and see what is happening. After applying the fix, we are still facing errors because *the server directive* doesn't exist on its own in .

Let's tell Lens Prism that, and see if it comes up with a better solution:

The server directive is not allowed in nginx.conf

You are correct. In a default nginx container, the main nginx.conf typically includes http, events, and other global directives. The server block is usually placed inside the http block, not at the top level.

If your nginx-config ConfigMap is being mounted as /etc/nginx/nginx.conf, it should look like this:

```
nginx
events {}

http {
    server {
        listen 80;
        location / {
            return 200 'Hello from the Lens tutorial!\n';
            add_header Content-Type text/plain;
        }
    }
}
```

 Copy

## Root Cause:

Your current nginx.conf likely places the server block at the top level, which is invalid. This causes nginx to fail to start, resulting in CrashLoopBackOff.

## Fix:

Wrap the server block inside an http block as shown above. Update the ConfigMap and restart the nginx pods.

Please verify and apply this change in Lens. Let me know if you need the exact kubectl patch command for the ConfigMap.

After applying this fix and restarting my nginx deployment, I see that all of my pods are running successfully:

default		Aa <input type="text"/> Search Pods...	<input type="button"/>	6 items
<input type="checkbox"/>	Name	<input type="button"/>	Namespace	Containers
<input type="checkbox"/>	hello-world	<input type="button"/>	default	<input type="checkbox"/>
<input type="checkbox"/>	instant-crash-pod	<input type="button"/>	default	<input type="checkbox"/>
<input type="checkbox"/>	jenkins-1754476299-fc8c98954-pl2gk	<input type="button"/>	default	<input type="checkbox"/>
<input type="checkbox"/>	nginx-deployment-5d5cc6d475-5k8bc	<input type="button"/>	default	<input type="checkbox"/>
<input type="checkbox"/>	nginx-deployment-5d5cc6d475-6ljpj	<input type="button"/>	default	<input type="checkbox"/>
<input type="checkbox"/>	nginx-deployment-5d5cc6d475-smssk	<input type="button"/>	default	<input type="checkbox"/>

Lens Prism sped up our debugging time considerably, and using at scale will help us solve many issues our Kubernetes resources will face. Because it is integrating with LLMs, you should always be mindful that fixes it provides may not always work on the first try, and that's to be expected.

Whenever a fix doesn't work, continue chatting with Lens Prism, because it supports multi-turn conversational queries, meaning that you can also ask follow-up questions, making your debugging process easier.

*Note: Lens Prism is available starting from the Plus/Pro tiers. To learn more about our pricing click [here](#).*

For more information about Lens Prism, check out this [article](#).

## Key points

In this article we've:

- Went through the most common K8s errors
- CrashLoopBackoff
- ErrImagePull/ImagePullBackOff
- CreateContainerError / CreateContainerConfigError
- OOMKilled
- NodeNotReady
- Unschedulable Pods: FailedScheduling
- Readiness/Liveness Probe Failures
- PodInitializing Error
- Saw the many ways we can solve these issues
- Went through how Lens K8s IDE speeds up the debugging process
- Understood how Lens Prism takes debugging to the next level

Knowing what to do when you are facing the most common Kubernetes errors saves you a lot of time, but as you saw throughout this article, there are many possible solutions to them.

Leveraging Lens K8s IDE and Lens Prism will help you a lot with your debugging process, and will equip you with all the tools you need for ensuring that your workloads are running smoothly inside of Kubernetes.

If you want to see Lens K8s IDE and Lens Prism in action, [download it today](#).

*Originally published at <https://k8slens.dev>.*

Kubernetes

DevOps

Kubernetes Cluster

OpenAI

AWS



## Published in k8slens

541 followers · Last published 4 days ago

Follow

Lens is the only IDE you'll ever need to take control of your Kubernetes clusters, and any applications and services running on top. It's built on open source and free. This is the official blog of Lens community.



## Written by Flavius Dinu

2.8K followers · 17 following

Follow



Developer Advocate | DevOps / Cloud Engineer | OpenTofu Evangelist |  
Docker Captain | AWS Community Builder  
<https://www.youtube.com/@devopswithflavius>

## No responses yet



David B Chase

What are your thoughts?

## More from Flavius Dinu and k8slens

**KUBERNETES DEPLOYMENT STRATEGIES**  
BLUE-GREEN, CANARY, ROLLING, A/B, SHADOW & MORE WITH REAL EXAMPLES

In k8slens by Flavius Dinu

### Mastering Kubernetes Deployment Strategies: Step-by-Step Guide...

Kubernetes offers many deployment options for your applications, and what options you...

Oct 13 16



...

Sep 29 19



...

**TOP 15 KUBERNETES USE CASES**

In k8slens by Flavius Dinu

### Top 15 Kubernetes Use Cases

Kubernetes has changed how organizations deploy, test, and scale their applications....



In k8slens by Flavius Dinu

## How to Set Up Prometheus and Grafana in Kubernetes with...

Monitoring your Kubernetes clusters is a mandatory step, but the process doesn't hav...

4d ago

3



...



Flavius Dinu

## Self-Host n8n with Docker: Should You Do It?

Let's be honest: Everyone is using n8n for automating their workflows. And I'm not goin...

Jul 8

64

2




...

See all from Flavius Dinu

See all from k8slens

## Recommended from Medium





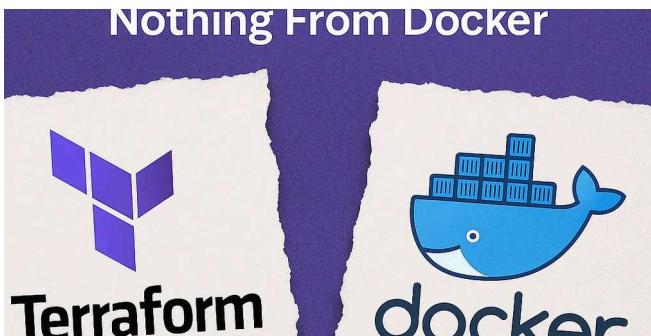
## 12 Observability Smells That Predict Outages

How to spot fragile telemetry, fix blind spots, and keep incidents from graduating into...

Oct 17 · 13



...



In Stackademic by Heinancabouly

## Terraform's License Drama Proves We've Learned Nothing From...

When will we stop building critical infrastructure on companies that can chang...

Oct 8 · 61



...



# Hash Block

## 12 Kubernetes Probes That Prevent Brownouts

Practical, production-ready health and readiness checks that keep traffic off...


## From CI/CD to AI/CD: The Next Evolution of DevOps in 2025

How Machine Intelligence Is About to Rewrite Your Deployment Pipeline—Whether You Li...

Oct 13 · 31



...



In DevOpsDynamo


## Terraform Done Right: How to Structure Projects That Scale

If you've ever opened an old Terraform repo and immediately regretted your life choices,...

Oct 17 · 19



...



In Dev Genius by Armin Nikdel Kourkah

## HULL: Simplifying Kubernetes Helm Charts

Helm has become the de facto package manager for Kubernetes. It allows developer...

⭐ Sep 24 ⌗ 11 🎙 2

↪+ ⋮ Sep 26 ⌗ 13

↪+ ⋮

See more recommendations